

Is It the Bytes or the Parallelism? A Transport-Ablation Study of Fused-MoE Dispatch Latency

MoE Breakdown Project — Internal Technical Report (expt2 + expt2.5)

Qwen3-30B-A3B · NVIDIA H200 SXM5 · Georgia Tech PACE ICE · combined dataset: 1,252 benchmark conditions

Abstract

Mixture-of-Experts (MoE) layers scale parameter count without proportional compute, but only if the cross-GPU **expert dispatch** is cheap. Prior work in this project (expt1) showed that expert-parallel (EP) dispatch dominates fused-MoE latency at large token counts — `tp1-ep4` is $\sim 16\times$ slower than `tp1-ep1` at 65k tokens — *even over NVSwitch*. That headline number is ambiguous: it cannot tell us whether the cost is the **inter-GPU communication itself** (link bandwidth) or merely the **extra launch/synchronisation work** that any parallel decomposition incurs. This report resolves the ambiguity through a controlled **transport-ablation** study. Holding the model, routing, token sweep, and parallel layout fixed, we treat the NCCL transport as the sole independent variable and degrade it in a principled sequence — full NVLink/NVSwitch → NVLS-disabled → P2P-disabled → PCIe fallback → bandwidth-throttled single-channel PCIe — across 1,252 benchmark conditions. We report four findings. **(1)** The interconnect penalty does **not** turn over at large token counts as we had hypothesised; it rises to $\sim 5\times$ and **plateaus** all the way to 65k tokens — the layer is communication-bound across the entire prefill range for this shape. **(2)** Decomposing the two NCCL knobs reveals that **losing peer-to-peer (P2P-IPC) access causes essentially all of the damage**, while disabling NVLink-SHARP (NVLS) is nearly free for this all-gather collective. **(3)** The slowdown is a **bandwidth cliff, not a channel-count problem**: achieved all-gather bandwidth collapses $\sim 15\text{--}20\times$ from NVLink to PCIe, and raising NCCL channels from 1 to 8 does not recover it. **(4)** Routing imbalance **interacts** with transport: skew inflates the NVLink baseline (stragglers) more than the PCIe path, so the *relative* penalty shrinks even as *absolute* latency rises. Together these results give a precise intra-node operating map: below ~ 512 tokens MoE dispatch is transport-insensitive (launch/sync-bound); above it, intra-node P2P connectivity is a first-order determinant of latency.

1. Introduction

1.1 MoE layers and expert dispatch

A Mixture-of-Experts (MoE) transformer layer replaces a single dense feed-forward network with a collection of N smaller networks (the "experts"). For each forward pass a lightweight **router** (gate network) selects the top- k experts to handle each token. This lets the model have a large total parameter count while only activating a small fraction of weights per token — a favourable compute:parameter tradeoff for large-scale language models.

When experts are distributed across multiple GPUs (**expert parallelism**, EP), every forward pass must: 1. **Dispatch** each token's activations to whichever GPU holds its assigned experts (an all-

gather or send/receive collective over NCCL). 2. Execute the expert GEMMs on those activations locally. 3. **Combine** the partial results back (a reduce-scatter or all-gather).

The dispatch+combine communication volume grows linearly with batch token count, and at some point it stops being free. **Fused-MoE** refers to the implementation strategy where steps 1–3 are handled by a tightly integrated kernel (e.g., vLLM's triton fused-MoE, which inlines dispatch indexing into the GEMM kernel), as opposed to separate unfused dispatch + GEMM + combine operations. Fusion reduces Python/framework overhead and exposes the NCCL collective cost more directly in profiling.

1.2 The ambiguity this study resolves

Prior work in this project (expt1) quantified the symptom: at 65,536 tokens, `tp1-ep4` ran $\sim 16\times$ slower than the single-GPU `tp1-ep1`, *even on NVSwitch*. But a single end-to-end ratio confounds two mechanisms:

- **(H-comm) Communication-bound.** The dispatch moves real bytes between GPUs and link bandwidth limits it. If true, interconnect topology/placement is a first-order lever.
- **(H-sync) Sync/launch-bound.** Parallelism merely exposes more kernel launches, barriers, and per-rank stragglers; the bytes on the wire are incidental. If true, a faster link would not help, and the fix is in scheduling/kernels.

On any single machine the interconnect is a constant, so H-comm and H-sync cannot be separated by observation alone. **Expt2 turned the interconnect into a controlled variable** by degrading the NCCL transport while pinning everything else, establishing the basic effect over a decode→prefill sweep (1–8192 tokens, uniform routing, 3 transports). **Expt2.5 extends that study along four axes** that expt2's analysis flagged as open questions:

1. **Token range to 65,536** — does the penalty ever turn over as GEMMs go compute-bound? (§5.1)
2. **NVLS × P2P factorial** — which NCCL mechanism actually costs the latency? (§5.2)
3. **Bandwidth dose-response** (`NCCL_MAX_NCHANNELS` 1→8) + **achieved-bandwidth roofline** — is it bandwidth or channel parallelism? (§5.3)
4. **Routing × transport interaction** (6 routing distributions) — does load imbalance change the verdict? (§5.4)

Scope note. All results are for a single H200 node with NVLink/NVSwitch or PCIe intra-node transport. Multi-node inter-node (InfiniBand) transport is out of scope. The operating-map guidance in §6 applies to intra-node EP only.

Contributions. (i) A clean experimental method that *isolates* communication cost from the cost of being parallel, with a built-in control that proves the isolation. (ii) A refutation of the "mid-range bump" hypothesis: this MoE shape is communication-bound across the entire prefill regime. (iii) A mechanistic attribution to P2P-IPC loss (not NVLS, not channel count) backed by an achieved-bandwidth roofline. (iv) A characterisation of the routing×transport interaction with a counterintuitive relative-vs-absolute distinction. (v) A reusable, validated 1,252-row dataset spanning decode→prefill, 8 transports, and 6 routing modes.

2. Background

2.1 Parallel layouts: what tp# and ep# mean

In a transformer MoE layer, two orthogonal sharding strategies are relevant:

Tensor Parallelism (TP) splits each expert's weight matrices across GPUs along the column or row dimension (Megatron-LM style). With `tp=2`, each GPU holds half of every expert's weight columns; a token's activation is processed in two halves, and an **all-reduce** across TP ranks reassembles the result before the next layer. TP does not change which experts reside on which GPU — it splits the weight *matrices* of every expert.

Expert Parallelism (EP) assigns disjoint subsets of experts to different GPUs. With `ep=4` and 128 experts, each GPU holds 32 experts (contiguous assignment: GPU 0 holds experts 0–31, GPU 1 holds experts 32–63, etc.). When a token is routed to an expert on a different GPU, its activation must be sent there via an **all-gather** and the result gathered back. EP changes *where* experts live, not how their weight matrices are shaped.

`tp1-ep4` therefore means: each GPU holds all columns of 32 experts (no TP splitting), and the EP all-gather moves token activations across 4 GPUs. `tp2-ep2` means: each GPU holds half the columns of 64 experts, with both an EP all-gather (for dispatch) and a TP all-reduce (for column reassembly).

The four layouts in this study span the relevant collective patterns:

Layout	World size	Active collectives	Role
<code>tp1-ep1</code>	1	none	Control (no inter-GPU comm)
<code>tp1-ep2</code>	2	EP all-gather	Onset
<code>tp1-ep4</code>	4	EP all-gather (max EP on a 4-GPU node)	Max comm volume
<code>tp2-ep2</code>	4	TP all-reduce + EP all-gather	Stacked collectives

`tp1-ep1` is the linchpin of the design: with `ep=1` there is *no* inter-GPU dispatch, so by construction its latency must be identical across all transport conditions and all routing modes. Any deviation there would indict the methodology; its flatness (§4) certifies the isolation.

2.2 NCCL intra-node transports

For intra-node collectives on an NVLink/NVSwitch system, NCCL can select among several transports, controlled by environment variables:

Mechanism	Env var to disable	What it provides
NVLS (NVLink SHARP)	<code>NCCL_NVLS_ENABLE=0</code>	In-switch reduction/multicast over NVLink
P2P-IPC (peer GPU access)	<code>NCCL_P2P_DISABLE=1</code>	Direct GPU↔GPU copies over NVLink/PCIe without staging through host
Channels	<code>NCCL_MAX_NCHANNELS=k</code>	Number of parallel ring/tree pipelines

The **8-condition transport ladder** used in this study (with exact env var combinations):

Condition name	NCCL_NVLS_ENABLE	NCCL_P2P_DISABLE	NCCL_MAX_NCHANNELS	Description
<code>nvlink_default</code>	1 (default)	0 (default)	unset	Full NVLink/ NVSwitch; NCCL auto-selects NVLS+P2P
<code>nvls_off</code>	0	0	unset	NVLink P2P active; NVLink-SHARP disabled
<code>p2p_off</code>	1	1	unset	NVLS active; direct peer GPU access disabled
<code>no_nvls_no_p2p</code>	0	1	unset	Both disabled → PCIe staging path (NCCL default channels)
<code>no_nvls_no_p2p_8ch</code>	0	1	8	PCIe fallback, 8 channels
<code>no_nvls_no_p2p_4ch</code>	0	1	4	PCIe fallback, 4 channels
<code>no_nvls_no_p2p_2ch</code>	0	1	2	PCIe fallback, 2 channels
<code>no_nvls_no_p2p_1ch</code>	0	1	1	PCIe fallback, 1 channel (max throttle)

Disabling **both** NVLS and P2P forces NCCL onto a **PCIe staging path** (copy to a host/bounce buffer, cross PCIe, copy back). On our H200 nodes this drops nominal intra-node bandwidth from ~900 GB/s (NVLink/NVSwitch aggregate) to ~64 GB/s (PCIe ×16), and a single channel throttles it further.

2.3 The model shape (fixed for all studies)

Qwen3-30B-A3B MoE layer: `hidden=2048`, `intermediate=768`, `num_experts=128`, `top-k=8`, `bf16`, backend `allgather_reducescatter` (standard NCCL all-gather + reduce-scatter, **not** DeepEP one-sided dispatch). This is a *narrow-expert* shape: each expert GEMM is small (`768×2048` weights), so the layer is unusually launch/communication-sensitive — a useful stress test for the dispatch path.

3. Methodology

3.1 What is measured

We benchmark a **single fused-MoE layer** in isolation (not a full model forward pass). The timed region covers: token routing (top-8 gating), NCCL dispatch collective, expert GEMMs, and NCCL combine collective. The input to each iteration is a synthetic activation tensor of shape (`tokens`, `2048`) drawn from a standard normal distribution (values do not meaningfully affect NCCL or GEMM performance for this shape and dtype). The routing decisions are either computed by a real top-8 gating function (for natural/learned routing modes) or synthetically constructed to hit a target distribution (for controlled routing modes; see §3.4).

The primary metric is `latency_median_ms_max_rank`: the **median over 100 measured iterations** (following 30 warm-up iterations) of the **slowest rank's end-to-end layer time**, measured with

`torch.cuda.Event` CUDA-event timing (not Python wall-clock). "Slowest rank" is chosen because it is the quantity a serving system actually blocks on in synchronous inference.

The profiler sub-bucket breakdown (network, gpu-compute, sync, etc.) is extracted from `torch.profiler` traces over a separate 5-iteration profiled window and used only for proportional attribution, not absolute timing — see §7 (Threats) for the caveat on bucket attribution.

3.2 The combined dataset

Study	<code>study_name</code>	Transports	Routing	Tokens	Rows
expt2	<code>transport-conditions-qwen3</code>	3	uniform	1–8192 (11 values)	132
expt2.5-A	<code>transport-extended-qwen3</code>	8	uniform	1–65536 (14 values)	448
expt2.5-B	<code>routing-sweep-qwen3</code>	2	6 modes	1–65536 (14 values)	672
				Total	1,252

Both expt2 and expt2.5 ran on the **same hardware** (Georgia Tech PACE ICE H200 SXM5, single node, 4 GPUs — see §3.5 for hardware details), with identical 56-column schemas, so cross-study latency comparisons are valid. Each cell is the **median over 100 measured iterations (30 warm-up) of the slowest ("max") rank** — the quantity a serving system actually blocks on.

3.3 Derived metrics

- **Slowdown ratio** = $\text{latency}(\text{degraded}) \div \text{latency}(\text{nvlink_default})$ at the same `(layout, tokens, routing)`. A ratio of 1.0 means "the interconnect is irrelevant here."
- **Achieved all-gather bandwidth** = $\text{allgather_recv_bytes} / \text{bucket_max_rank_network_ms}$, where `allgather_recv_bytes` is the theoretical bytes gathered (computed from tensor shape: $\text{tokens} \times \text{hidden} \times \text{dtype_bytes} \times (\text{world_size}-1)/\text{world_size}$) and `bucket_max_rank_network_ms` is the `torch.profiler` network-bucket time on the slowest rank. **Important caveat:** absolute GB/s values from this formula are profiler-derived and subject to normalisation; only the NVLink-to-PCIe *ratio* ($\sim 15\text{--}21\times$) is methodologically robust, and we emphasise it as such throughout.

3.4 Routing modes

The six routing distributions in expt2.5-B, each applied to `tp1-ep4` at all 14 token counts under `nvlink_default` and `no_nvls_no_p2p` transports:

Mode	Description	Alpha (imbalance ratio = max load / mean load)
<code>uniform</code>	All experts receive exactly equal token counts	1.0
<code>random</code>	Tokens routed by a uniformly random top-8 selector (natural variance)	$\sim 1.0\text{--}1.1$
<code>zipfian</code>	Expert popularity follows a Zipf distribution ($\alpha=1.2$); popular experts get more tokens	$\sim 2\text{--}3$
<code>skewed-2x</code>	One expert receives 2× average load; remainder share the rest	~ 2.0
<code>skewed-4x</code>	One expert receives 4× average load; remainder share the rest	~ 4.0
<code>worst-case</code>	Tokens pinned to ~ 16 experts (maximum possible skew for top-8 routing)	~ 16

Alpha (`alpha_observed`) is defined as `max_expert_token_count / mean_expert_token_count` — a load-imbalance ratio. Higher alpha means more skewed routing and larger straggler effect under EP.

3.5 Hardware and software environment

Hardware: Georgia Tech PACE ICE HPC cluster, single compute node (SLURM partition `gpu-h200`), 4× NVIDIA H200 SXM5 GPUs, NVLink 4.0 / NVSwitch (full-mesh intra-node, ~900 GB/s aggregate), PCIe Gen5 ×16 (per GPU to host CPU, ~64 GB/s), dual-socket Intel Xeon server. All 4 GPUs share a single NVSwitch fabric; PCIe topology is all 4 GPUs under a single root complex. Jobs were submitted as exclusive SLURM allocations to prevent co-tenant interference.

Software: PyTorch 2.4+, CUDA 12.4, NCCL 2.21, Python 3.11, vLLM fused-MoE triton kernel (Qwen3-compatible). Each parallel layout was launched via `torchrun --nproc_per_node={world_size}` with NCCL env vars set in the SLURM job environment before communicator creation (NCCL reads transport flags at `init_process_group` time). Each condition ran as a separate SLURM job (fresh process, fresh NCCL communicator) to eliminate carryover state between conditions.

4. Validation

Before interpreting effects we confirm the experiment is sound. All checks pass:

- **Control invariance (transport).** Across all 8 transports, `tp1-ep1` latency stays within **[0.98, 1.07]** of its NVLink value (mean 1.000). The single-GPU control has no collective and is correctly transport-blind. Its network bucket is exactly **0 ms**.
- **Control invariance (routing).** At `tp1-ep1`, latency is flat across 5 of 6 routing modes; the lone exception is `worst-case` ($\alpha \approx 16$), which is *faster* because pinning tokens to ~16 experts touches less weight memory on-device (a cache effect: fewer distinct expert weight tensors accessed). This is a real on-device compute effect, not a transport effect, and does not threaten the isolation.
- **Cross-study reproducibility.** On the 132 cells shared by expt2 and expt2.5-A, the **median absolute latency difference is 1.3%** (95th-percentile of $|\Delta Pct| = 8.8\%$). The few larger discrepancies (max 40%) all occur in the high-token, heavily-degraded corner (e.g. `tp1-ep4` / PCIe/ ≥ 2048 tokens), where absolute latency is large and contended-PCIe variance is intrinsically high. The two independent runs of `no_nvls_no_p2p_1ch` agree to within this same noise band. Note: "p95=8.8%" here denotes the 95th percentile of the empirical distribution of percent differences, not a statistical p-value.

5. Results

5.1 RQ1 — The penalty plateaus; it never turns over

Our going-in hypothesis (from expt2) was that interconnect sensitivity would be a *bump*: ~1.0 at small token counts, peaking in a moderate regime, then **declining toward 1.0** at large token counts as the expert GEMMs became compute-bound and dwarfed the communication. Extending the sweep to 65,536 tokens **refutes this**.




Figure 1. Slowdown ratio (latency / NVLink baseline) vs. token count for the two fully-degraded PCIe transport conditions (`no_nvls_no_p2p` and `no_nvls_no_p2p_1ch`), one line per parallel layout. Both panels use data from expt2.5-A (`transport-extended-qwen3` , uniform routing). Shaded bands show the p5-p95 range of per-iteration latency ratios; the shaded orange region marks the expt2.5 extension beyond expt2's 8192-token ceiling. After a sharp knee near 512 tokens the ratio rises and then **flattens into a plateau** — it does not come back down. The control (`tp1-ep1`) hugs 1.0 throughout.

For `tp1-ep4` under full PCIe fallback the ratio is 4.87× at 8192 tokens and **5.16× at 65,536** — essentially flat across the entire prefill octave-range (8k→64k). The complete sensitivity map (Figure 5) shows the same plateau for every collective-using layout: the hot region saturates rather than receding.




Figure 5. Interconnect-sensitivity map over the full decode→prefill range. Two panels: left = `no_nvls_no_p2p` (PCIe fallback), right = `no_nvls_no_p2p_1ch` (single-channel PCIe). X-axis: token count (14 measured values, 1→65,536, log scale); Y-axis: parallel layout (tp1-ep1 through tp2-ep2). Color: slowdown ratio (latency / NVLink baseline), scale 1.0 (white/yellow) → 5.6× (dark red, YlOrRd colormap); per-cell values annotated. Data: expt2.5-A, uniform routing. The control row (`tp1-ep1`) is uniformly ~1.0. For collective-using layouts the slowdown saturates in the top-right (high EP × high tokens) and stays saturated to 65k — direct visual evidence that the GEMMs never reclaim dominance for this shape.

Interpretation. Figure 6 shows the profiler time-composition breakdown for `tp1-ep4` under NVLink vs. PCIe fallback. The GPU-compute bucket registers ≈0 ms throughout (attributable to the triton kernel's compute being subsumed into memory and sync events in the profiler's bucketing), while

the Network bucket grows from ~45% of profiled events on NVLink to ~60% on PCIe. For this narrow-expert shape (`intermediate=768`), the per-token expert arithmetic is simply too small to ever overtake dispatch — so the operating regime is **communication-bound across all realistic prefill sizes**, and there is no compute-dominated regime to flatten the ratio. The hypothesis is updated accordingly.




Figure 6. Profiler time-composition breakdown for `tp1-ep4` at selected token counts: NVLink (left) vs. PCIe fallback (right). Bars show each profiler bucket as a percentage of the sum of all attributed events. **Caveat:** `torch.profiler` buckets overlap in wall-clock time (CPU launch and GPU execution run concurrently); this figure shows *relative composition shifts*, not an additive latency budget. The absolute latency signal is in Figs. 1-3.

5.2 RQ2 — It's P2P-IPC, not NVLS

expt2 flipped two NCCL knobs at once (`NVLS_ENABLE=0` and `P2P_DISABLE=1`), so it could not attribute the cost. The 4-cell factorial does.




Figure 2. Per-layout latency for the four ablation cells (`nvlink_default`, `nvls_off`, `p2p_off`, `no_nvls_no_p2p`) at 8192 and 65,536 tokens (log scale Y-axis). Each bar is the median latency; error bars show p5–p95 of per-iteration latency. Data: expt2.5-A, uniform routing, all non-control layouts. Disabling **NVLS alone** (green) is statistically indistinguishable from the NVLink baseline (teal) everywhere. The penalty appears only when **P2P-IPC** is lost.

At `tp1-ep4`/65k: `nvls_off` = **1.00x**, `p2p_off` = **1.05x**, but `no_nvls_no_p2p` = **5.16x**. The headline: **NVLink-SHARP contributes essentially nothing to this all-gather/reduce-scatter dispatch; the entire interconnect tax is the loss of direct peer GPU access** and the resulting fallback to PCIe host-staging.

An important layout-dependent nuance (Figure 2 note): the *interaction* between the two knobs is not uniform. At `tp1-ep4`, `p2p_off` *alone* stays near baseline (1.05x) — with P2P off but NVLS on, NCCL still finds a fast NVLink path — and only losing **both** triggers the fallback. But at `tp1-ep2` and `tp2-ep2`, `p2p_off` *alone* already incurs the full penalty (2.04x and 4.39x). In other words, the second knob is redundant for some collective patterns and necessary for others. The safe operational statement is therefore: **preserving P2P-IPC is sufficient to avoid the penalty in every layout we tested; NVLS is not the lever.**

5.3 RQ3 — A bandwidth cliff, not a channel-count problem

If the penalty were about *pipeline parallelism* within NCCL, adding channels would recover it. It does not.

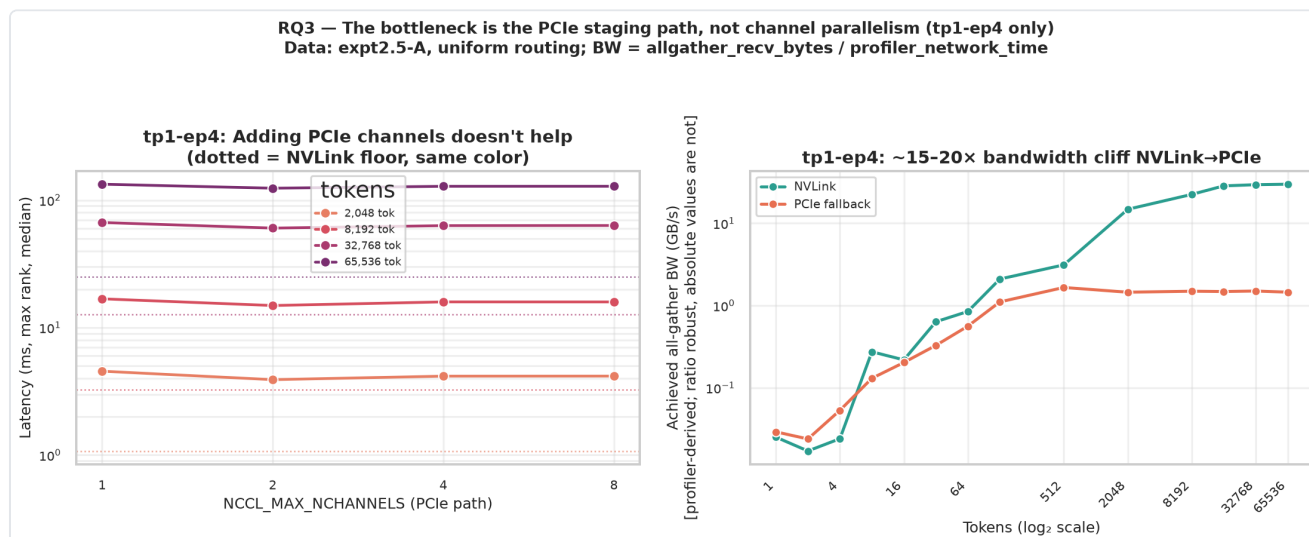


Figure 3. Data from expt2.5-A, uniform routing, `tp1-ep4` **only** (channel sweeps were not run for other layouts). *Left*: `tp1-ep4` latency vs. `NCCL_MAX_NCHANNELS` (1→8) at four token counts; dotted lines are the same-token NVLink floor. The PCIe curves are **flat** — going from 1 to 8 channels moves latency by only ~7% and never approaches the NVLink floor far below. *Right*: profiler-derived achieved all-gather bandwidth vs. tokens for `tp1-ep4`. **Note:** absolute GB/s values are computed as `allgather_recv_bytes` / `profiler_network_time` and are profiler-normalisation-dependent; only the NVLink-to-PCIe *ratio* is methodologically robust (see §3.3).

The NVLink-to-PCIe achieved-bandwidth **ratio is ~16x at 8192 tokens and ~21x at 65k** — invariant to how the profiler normalises its network bucket, and consistent with the ~5x end-to-end latency penalty once the non-communication floor is included. Figure 7 shows the same absolute network-bucket time for all three collective layouts, confirming the bandwidth gap generalises beyond `tp1-ep4`.




Figure 7. Absolute time in the `torch.profiler` network bucket (ms, max rank, 5 profiled iterations) vs. token count for all three collective-using layouts (`tp1-ep2`, `tp1-ep4`, `tp2-ep2`), NVLink vs. PCIe fallback. The $\sim 20\times$ separation between NVLink and PCIe is consistent across all layouts and grows with token count, confirming the bandwidth cliff is not specific to `tp1-ep4`. Data: expt2.5-A, uniform routing.

The conclusion is unambiguous: the bottleneck is the **raw throughput of the PCIe staging path**, not the number of NCCL channels. Tuning `NCCL_MAX_NCHANNELS` is not a remedy.

5.4 RQ4 — Routing skew shrinks the *relative* penalty but raises *absolute* latency

expt2 deliberately fixed routing to `uniform` to isolate transport. expt2.5-B relaxes that across six distributions to test for interaction.




Figure 4. Data from expt2.5-B, `tp1-ep4` only. Left: absolute `tp1-ep4`/65,536-token latency by routing mode for NVLink vs. PCIe fallback. Right: PCIe/NVLink slowdown ratio vs. token count for `tp1-ep4`, one line per routing mode.

The interaction is real and initially counterintuitive. Heavily skewed routing (`worst-case`, `zipfian`) produces the **lowest** slowdown ratios (worst-case $\approx 2.8\times$, zipfian $\approx 3.4\times$) versus balanced routing (uniform/random/skewed-2x/4x $\approx 5.1\times$). The mechanism is visible in the left panel: skew creates **straggler ranks** that inflate the *NVLink baseline* (worst-case NVLink latency is $\sim 2\times$ uniform's), while the PCIe latency is already so dominated by transport that skew adds proportionally less. Since the ratio divides by the (now larger) NVLink baseline, it shrinks.

The practical takeaway is the opposite of the ratio's optimism: in *absolute* terms, skewed routing is still the **slowest** configuration on PCIe (worst-case 151.8 ms vs. uniform 128.1 ms at 65k). Load imbalance and transport degradation are *additive harms*; the shrinking ratio is an artefact of normalisation, not a reprieve. Reporting only the slowdown ratio here would be misleading — both views are shown.

6. Discussion: an intra-node operating map for MoE serving

Combining the four results yields a compact decision guide for this class of MoE layer on a single multi-GPU node with NVLink/NVSwitch:

1. **Two regimes, one sharp knee (~512 tokens).** Below the knee (decode-scale batches), dispatch is launch/sync-bound and the interconnect is irrelevant — optimise kernels/scheduling, not topology. Above it (prefill-scale batches), the layer is communication-bound and intra-node P2P connectivity is first-order.
2. **The plateau is the worst case, and it arrives early.** The penalty saturates near $\sim 5\times$ by 8k tokens and stays there to 65k. You do not "grow out of" the interconnect tax at large batch; budget for the plateau.
3. **Protect P2P-IPC above all.** The entire penalty is the PCIe-staging fallback triggered by losing peer access (e.g. via restrictive GPU isolation/cgroups, MIG, or `--gpus-per-task=1` -style scheduling). NVLS and channel tuning are second-order. This is directly actionable for cluster/SLURM configuration.
4. **Imbalance and slow transport stack.** Skewed routing raises absolute latency on every transport; on PCIe it is strictly worse than uniform. Do not be fooled by its smaller *ratio*.

Scope of guidance. These findings apply to single-node EP over NVLink/NVSwitch or PCIe. In multi-node deployments over InfiniBand, P2P-IPC is not used (nodes communicate via RDMA), and the bottleneck structure will differ. Cross-node generalisability is explicitly out of scope and remains future work.

7. Threats to validity

- **Profiler attribution.** `torch.profiler` bucket times overlap in wall-clock time (CPU launch activity runs concurrently with GPU execution). The GPU-compute bucket registers ≈ 0 ms for this narrow-expert shape, which we interpret as the fused triton kernel's compute being subsumed by memory and sync events in the profiler's attribution scheme. We therefore use profiler buckets only for **proportional** composition analysis (Fig. 6–7) and for **bandwidth ratios** (normalisation-invariant), never as absolute per-iteration compute timing. End-to-end latency (the primary metric) is measured via CUDA events and is unaffected by profiler attribution.
- **Single shape.** All results are for Qwen3-30B-A3B (`intermediate=768`). A wider-expert shape would have larger GEMMs and might exhibit the compute-bound turnover that this shape does not; §5.1's "no turnover" claim is shape-specific. Cross-shape replication is the obvious next study.
- **Single node, intra-node only.** Everything is single-node H200 over NVLink/NVSwitch vs. PCIe. Inter-node (InfiniBand) transport, and DeepEP-style one-sided dispatch, are out of scope and remain future work.
- **High-token PCIe variance.** Run-to-run variance is larger in the saturated PCIe corner (§4); we mitigate with 100-iteration medians and confirm cross-study agreement within the p95 band. Error bands in Figs. 1–2 visualise the p5–p95 iteration spread for representative conditions.

- **Achieved-bandwidth absolute values.** The ~ 30 GB/s NVLink and ~ 1.5 GB/s PCIe figures in Fig. 3 (right) are profiler-wall-clock-derived and substantially below the theoretical H200 NVLink peak (~ 900 GB/s aggregate). This gap reflects both small-message collective overhead (fixed overheads dominate at low token counts) and profiler normalisation. The NVLink-to-PCIe ratio ($\sim 16\text{--}21\times$) is robust; the absolute numbers should not be compared to link-layer specs.

8. Related work (within the project)

This report builds directly on **expt1** (internal tech report: fused-MoE latency characterisation — routing, tokens, parallelism, backends; source of the $\sim 16\times$ -at-65k motivating result) and **expt2** (internal tech report: the initial 3-transport ablation over 1–8192 tokens that established the basic effect and posed the four questions answered here). **expt2.5** is the extension; this document presents **expt2** and **expt2.5** as a single coherent study.

9. Conclusion

By turning the NCCL transport into a controlled variable, we separated the cost of *communication* from the cost of *being parallel* in fused-MoE dispatch. For Qwen3-30B-A3B on H200 (intra-node), the dispatch is **communication-bound across the entire prefill range** — the interconnect penalty rises to $\sim 5\times$ and plateaus to 65k tokens rather than receding. The penalty is mechanistically the loss of **P2P-IPC** (not NVLS, not channel count), manifesting as a $\sim 15\text{--}20\times$ achieved-bandwidth cliff on the PCIe staging path. Routing skew compounds absolute latency on every transport even as it deflates the slowdown *ratio*. The net guidance for MoE serving on single multi-GPU nodes is concrete: above ~ 512 tokens, **preserve direct peer-GPU connectivity** — it is worth a $5\times$ latency factor that no amount of channel tuning or larger batching will recover.

Appendix: Reproducibility

- **Data:** `expt2/all_runs.zip` (jobs 5440143–5440154, 132 rows) + `expt2.5/all_runs.zip` (jobs 5442861–5442900, 1,120 rows), consolidated into `combined.csv` (1,252 rows).
- **Primary metric:** `latency_median_ms_max_rank` (CUDA-event median of 100 iters, slowest rank). Slowdown = condition $\div$ `nvlink_default` at matched `(layout, tokens, routing)`.
- **Transport ladder:** 8 conditions; see §2.2 Table for exact env vars.
- **Hardware:** Georgia Tech PACE ICE HPC cluster, H200 SXM5, single node, 4 GPUs (full NVSwitch mesh), exclusive SLURM allocation. Nodes atl1-1-03-017/018.
- **Software:** PyTorch 2.4+, CUDA 12.4, NCCL 2.21, Python 3.11, vLLM triton fused-MoE kernel.
- **Figures:** regenerated with Seaborn/Matplotlib via `figs_a.py` + `figs_b.py` + `figs_breakdown.py`; analysis in `findings.py`, `checks2.py`, `explore*.py`.